

TÁMBIA FURRER

GPAI Compliance: Framework, Regulatory Matrix & Operational Tool

Supporting materials for the AI Compliance Officer application

A Framework for Operationalizing GPAI Compliance

The 8-Stage Compliance Lifecycle | Executive Briefing

GPAI providers face a regulatory landscape that has no clean precedent. The EU AI Act’s Articles 51–56 create obligations that operate at model scale rather than system scale, flow downstream to deployers through structured information packages, and require demonstrable technical capability rather than procedural attestation. Layered atop the AI Act are active training-data litigation, copyright and text-and-data-mining opt-out regimes, export controls on model weights, voluntary commitments to governments, and formal engagement with AI Safety Institutes. Most of these obligations did not exist three years ago. Several do not yet have settled interpretive guidance.

In this environment, the limiting factor for sustainable compliance is not regulatory knowledge alone but the operational infrastructure to apply that knowledge consistently across model development, release, monitoring, and downstream relationships. The eight-stage framework introduced in this briefing is designed to help GPAI leadership organize that infrastructure — to see clearly what compliance work the organization is doing, what work it is not doing, and how the pieces should connect.

Purpose of This Briefing

To introduce AI leadership to a structured framework for operationalizing AI regulatory obligations across the full compliance lifecycle — from initial discovery through continuous monitoring — and to make explicit how each stage maps to obligations under the EU AI Act, Codes of Practice, and adjacent legal regimes.

Why GPAI Provider Compliance Requires Its Own Framework

Compliance frameworks designed for AI deployers — or worse, retrofitted from privacy compliance — fail to capture what makes GPAI provider obligations structurally distinct. Six characteristics define the GPAI compliance posture.

- **Provider/deployer asymmetry.** Obligations flow one-to-many from the GPAI provider to many downstream deployers. The information packages a provider supplies determine whether deployers can fulfill their own obligations.
- **Systemic-risk tier dominance.** For models above the 10^{25} cumulative-training-FLOP threshold under Article 51, Article 55 imposes model evaluation, adversarial testing, systemic risk assessment, serious incident reporting, and cybersecurity obligations that are materially more demanding than the baseline.
- **Training data as foundation.** Training-data governance — provenance, licensing, opt-out compliance, and the Article 53 publicly available training-data summary — sits at the center of the compliance work, with no equivalent in deployer compliance.
- **Public-facing artifacts.** Training-data summaries, copyright policies, and model cards are public documents subject to regulatory and public scrutiny. They function as compliance artifacts and as market-facing communications simultaneously.
- **Technical capability is the obligation.** Article 55 requires demonstrable capability for model evaluations, adversarial testing, and dangerous-capability assessment. This is technical work, not documentation work.
- **Government safety body engagement.** Pre-deployment evaluation arrangements with AI Safety Institutes, voluntary safety commitments, and government testimony are operational realities, not occasional events.

The 8-Stage Compliance Lifecycle

The framework organizes GPAI compliance into eight sequential but continuously interacting stages. Each stage answers a distinct question that compliance leadership must be able to answer with evidence at any moment.

#	Stage	Representative Operations	Representative Artifacts
1	Discovery & Understanding	Inventory of models, training data, compute, downstream deployers, capabilities, and access tiers.	Model portfolio register; training-data inventory; cumulative-compute log; downstream deployer registry; capability-evaluation registry.
2	Assessment Against Requirements	Article 51, 53, and 55 analysis; Code of Practice analysis; copyright and adjacent law analysis.	Regulatory applicability memos; systemic-risk threshold assessment (Art. 51); Code of Practice gap analysis (Art. 56); legal opinion archive.
3	Gap Identification	Comparison of current practice against requirements, combining documentation review with technical capability assessment.	Gap assessment reports; technical capability gap analyses; senior-management findings memo; remediation prioritization document.
4	Decision-Making & Prioritization	Model release, capability disclosure, voluntary commitment, access tier, and Safety Institute engagement decisions.	Frontier Model Safety Committee minutes; model release approval records; risk acceptance memos; voluntary commitment register.

#	Stage	Representative Operations	Representative Artifacts
5	Operational Execution	Model evaluations, adversarial testing, training-data governance, downstream information provision, serious-incident response.	Evaluation reports; red-team reports; training-data governance procedures; downstream information packages (Art. 53(1)(b)); incident response runbooks.
6	Documentation & Evidence	Article 53 technical documentation, public artifacts, evaluation evidence, systemic risk assessments.	Technical documentation (Art. 53(1)(a), Annex XI); training-data summary (Art. 53(1)(d)); copyright policy (Art. 53(1)(c)); model and system cards; serious-incident records (Art. 55).
7	Demonstration & Communication	Publication of required artifacts, AI Office engagement, Safety Institute coordination, downstream and regulator response.	Published training-data summary; AI Office correspondence; AISI pre-deployment testing reports; downstream provider information packages; serious-incident reports (Art. 55(1)(c)).
8	Continuous Monitoring	Ongoing tracking of compute, capability evolution, misuse, legal landscape, and regulatory development.	Post-market monitoring reports; capability-evolution logs; regulatory horizon-scanning briefings; incident trend reports; quarterly compliance dashboards.

The stages are presented sequentially for clarity, but in practice they operate as a continuous loop. Stage 8 monitoring surfaces new findings that trigger return to Stage 1 discovery; Stage 4 decisions reshape Stage 5 execution priorities; Stage 7 demonstration reveals gaps that route back through Stages 3 and 4.

How the Framework Operationalizes Regulatory Obligations

The framework is not a checklist or a maturity model. It is an organizing structure that lets compliance leadership translate specific regulatory obligations into specific operational practices, with clear accountability for each.

From legal text to operational work. Each obligation under Articles 51–56, each commitment under an Article 56 Code of Practice, and each requirement under adjacent regimes (copyright, export control, voluntary commitments) maps to one or more activities within specific stages. The framework lets leadership see at a glance which activities address which obligations, where coverage is complete, where it is partial, and where it is absent.

From individual obligations to integrated programs. Many GPAI obligations interact. Training-data summary publication (Article 53) depends on training-data inventory (Stage 1) and copyright policy implementation (Stage 5). Systemic-risk assessment (Article 55) depends on capability evaluations (Stages 1 and 5) and adversarial testing (Stage 5), generating evidence for documentation (Stage 6) and disclosure (Stage 7). The framework reveals these dependencies and prevents the fragmentation that occurs when obligations are treated in isolation.

From point-in-time compliance to continuous compliance. GPAI obligations are not satisfied once and forgotten. Training compute accumulates and may cross the systemic-risk threshold. Models develop new capabilities post-training. Codes of Practice evolve. The framework’s Stage 8 makes continuous monitoring an explicit, structured activity rather than ad-hoc reaction to events.

From compliance posture to demonstrable evidence. Regulators, downstream deployers, customers, and courts increasingly demand evidence rather than attestation. Stages 6 and 7 of the framework structure the evidence base and the demonstration capability that lets the organization respond to any inquiry — from an AI Office investigation to a litigation discovery request — with documented, traceable, defensible material.

Using the Framework

For GPAI executives evaluating their current compliance posture, four uses of the framework are most immediately valuable.

- **Map current activities to stages.** Identify which activities the organization performs in each stage, where ownership sits, and where gaps exist. The mapping itself frequently reveals coverage gaps and ownership ambiguity that have gone unnoticed.
- **Trace obligations through the lifecycle.** For each major obligation — Article 53 baseline, Article 55 systemic risk, training-data summary publication, copyright policy, serious incident reporting, downstream information provision — confirm that supporting activities exist across all relevant stages.
- **Align cross-functional accountability.** GPAI compliance spans legal, policy, safety research, engineering, product, and executive functions. The framework supports a shared language for cross-functional coordination and clear accountability assignment by stage and activity.
- **Anticipate regulatory inquiry response.** When a regulator, deployer, or court asks for evidence of compliance with a specific obligation, the framework lets the organization respond with structured, comprehensive material rather than scrambling to assemble it under deadline pressure.

Closing Observation

The organizations that succeed in GPAI compliance over the next several years will not be those with the most sophisticated legal interpretations or the most advanced safety research, considered in isolation. They will be those that operationalize their compliance capability — building the durable infrastructure that lets legal interpretation, safety research, product decisions, and downstream relationships work together coherently as the regulatory landscape continues to develop.

The 8-Stage framework is offered as a structure for that work. It is not the only viable structure, and it will evolve as the regulatory landscape evolves. But it provides a starting point that translates the abstraction of “GPAI compliance” into the concrete operational practices that make compliance sustainable at scale. A companion document provides the detailed activity breakdown for each stage.

GPAI Provider Pre-Deployment Compliance Activities

Activity-by-Regulation Matrix with Second-Line Review and Sign-Off Designations

The matrix below maps each compliance activity to the regulatory sources that require it, and indicates whether the activity requires second-line review only (R) or both review and sign-off (R+S). The activities are organized into eleven operational categories spanning model lifecycle work, baseline GPAI documentation, transparency obligations, AI literacy, external engagement, adjacent regulatory regimes, and internal governance. Activities marked with checkmarks indicate the regulatory sources from which the obligation directly arises; compliance officers should verify each is complete or appropriately tracked before deployment authorization is recommended.

R	Second-line review required	R+S	Review and sign-off required	✓	Activity required by this regulatory source
SRM	Systemic Risk Model (Art. 51 threshold met or designated)	GPAI	Baseline GPAI model	PEND	Classification pending (approaching threshold)

Activity	Applies To	2nd Line	Art. 53	Art. 55	Art. 56	Art. 50	Art. 4	Art. 75	GDPR	Copy-right	Export Ctrl	Vol. Cmts
1. PRE-TRAINING												
Model Architecture Design	SRM GPAI PEND	R		✓	✓							
Training Planning	SRM GPAI PEND	R	✓	✓	✓							
Dataset Selection	SRM GPAI PEND	R+S	✓	✓	✓				✓	✓		
2. TRAINING												
Early Capability Evaluation	SRM PEND	R		✓	✓							
Safety Hypothesis	SRM PEND	R		✓	✓							
3. PRE-DEPLOYMENT EVALUATION												
Phase 2 Capability Evaluation	SRM PEND	R		✓	✓							
Adversarial Testing	SRM PEND	R+S		✓	✓							
Red-Teaming Evaluation	SRM PEND	R		✓	✓							
Safeguard Testing	SRM GPAI PEND	R+S		✓	✓	✓						
Deployment Simulations	SRM PEND	R		✓	✓							
AISI Pre-Deployment Evaluation	SRM PEND	R+S		✓	✓							✓

Activity	Applies To	2nd Line	Art. 53	Art. 55	Art. 56	Art. 50	Art. 4	Art. 75	GDPR	Copy-right	Export Ctrl	Vol. Cmts
4. PRE-DEPLOYMENT COMPLIANCE ASSESSMENT												
Cybersecurity Assessment	SRM PEND	R+S		✓	✓							
Systemic-Risk Analysis	SRM PEND	R+S		✓	✓							✓
5. DEPLOYMENT AUTHORIZATION												
Model Report	SRM GPAI PEND	R+S	✓	✓	✓							
Risk Assessment	SRM GPAI PEND	R+S		✓	✓				✓			
Mitigation Summary	SRM GPAI PEND	R		✓	✓							
Proposed Deployment Controls	SRM GPAI PEND	R+S	✓	✓	✓	✓						
Residual Risk Recommendation	SRM GPAI PEND	R		✓	✓							
6. BASELINE GPAI DOCUMENTATION												
Technical Documentation (Annex XI)	SRM GPAI PEND	R+S	✓		✓							
Downstream Provider Information Package	SRM GPAI PEND	R+S	✓		✓							
Copyright Policy	SRM GPAI PEND	R+S	✓		✓					✓		
Training Data Summary (Public)	SRM GPAI PEND	R+S	✓		✓					✓		
7. TRANSPARENCY & DISCLOSURE												
AI Interaction Disclosure Mechanisms	SRM GPAI PEND	R				✓						
AI-Generated Content Marking	SRM GPAI PEND	R+S				✓						

Activity	Applies To	2nd Line	Art. 53	Art. 55	Art. 56	Art. 50	Art. 4	Art. 75	GDPR	Copy-right	Export Ctrl	Vol. Cmts
Deepfake Labeling Mechanisms	SRM GPAI PEND	R+S				✓						
Automated Decision-Making Notices	SRM GPAI PEND	R				✓			✓			
8. AI LITERACY & TRAINING												
Staff AI Literacy Program	SRM GPAI PEND	R					✓					
Role-Specific AI Training Records	SRM GPAI PEND	R					✓					
Training Effectiveness Assessment	SRM GPAI PEND	R					✓					
9. EXTERNAL ENGAGEMENT												
AI Office Engagement Log	SRM GPAI PEND	R	✓	✓				✓				
National Competent Authority Engagement	SRM GPAI PEND	R		✓				✓				
Code of Practice Adherence Statement	SRM GPAI PEND	R+S			✓							
AISI / Safety Institute Engagement	SRM PEND	R+S		✓								✓
Serious Incident Reporting Procedures	SRM PEND	R+S		✓				✓	✓			
Conformity Assessment Body Coordination	SRM PEND	R+S		✓	✓							
10. ADJACENT REGULATORY REGIMES												
GDPR Legal Basis Assessment (training data)	SRM GPAI PEND	R+S							✓			
Data Protection Impact Assessment	SRM GPAI PEND	R+S							✓			
Cross-Border Transfer Mechanisms	SRM GPAI PEND	R+S							✓			

Activity	Applies To	2nd Line	Art. 53	Art. 55	Art. 56	Art. 50	Art. 4	Art. 75	GDPR	Copy-right	Export Ctrl	Vol. Cmts
TDM Opt-Out Compliance Mechanisms	SRM GPAI PEND	R+S	✓							✓		
Training Data Licensing Documentation	SRM GPAI PEND	R+S	✓							✓		
Export Control Classification	SRM GPAI PEND	R+S									✓	
Model Weight Distribution Controls	SRM PEND	R+S		✓							✓	
Voluntary Commitment Adherence (Hiroshima, etc.)	SRM GPAI PEND	R										✓
Industry Coalition Commitments (FMF, PAI)	SRM GPAI PEND	R										✓
Cross-Jurisdictional Compliance Assessment	SRM GPAI PEND	R+S	✓	✓					✓	✓	✓	✓
11. INTERNAL GOVERNANCE												
Quality Management System Documentation	SRM GPAI PEND	R+S	✓	✓	✓							
Frontier Model Safety Committee Sign-Off	SRM PEND	R+S		✓	✓							✓
Board / Governance Body Notification	SRM GPAI PEND	R+S	✓	✓								
Executive Risk Acceptance	SRM GPAI PEND	R+S		✓								
Legal Sign-Off on Compliance Posture	SRM GPAI PEND	R+S	✓	✓	✓	✓			✓	✓	✓	
Compliance Officer Pre-Deployment Attestation	SRM GPAI PEND	R+S	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓

Model classification: SRM (Systemic Risk Model) — a GPAI model presumed to present systemic risk under Article 51 because cumulative training compute exceeds 10²⁵ FLOPs, or designated as such by the AI Office on the basis of capabilities, reach, or other indicators. Subject to the full Article 55 obligation stack in addition to Article 53. **GPAI (Baseline GPAI)** — a general-purpose AI model below the systemic risk threshold and not otherwise designated. Subject to Article 53 baseline obligations but not Article 55 systemic risk obligations. **PEND (Pending Classification)** — a model approaching the Article 51 threshold or under active classification review. Treated as SRM-equivalent for compliance planning purposes to avoid a scramble if reclassification occurs.

Regulatory source key: Art. 53 (baseline GPAI) | Art. 55 (systemic-risk GPAI) | Art. 56 (Codes of Practice) | Art. 50 (transparency to natural persons) | Art. 4 (AI literacy) | Art. 75 (authority cooperation) | GDPR (Regulation (EU) 2016/679) | Copyright (DSM Directive Art. 4(3) and adjacent IP regimes) | Export Ctrl (EU Reg. 2021/821, US EAR, equivalent national regimes) | Vol. Cmts (Hiroshima Process, AI Safety Summit commitments, Frontier Model Forum, Partnership on AI).


Notes. Activity-to-regulation mapping reflects the source from which the obligation directly arises; many activities also produce evidence relevant to other regulatory sources. Second-line designations (R / R+S) reflect the compliance function's oversight role; first-line execution is owned by the Designated Responsible Team. Sectoral overlays (financial services, healthcare, employment) would add columns where the model is deployed into regulated sectors. Sub-jurisdictional regimes (UK, US state laws, APAC frameworks) are tracked under Cross-Jurisdictional Compliance Assessment.

Atlas — A Conceptual Tool for Operationalizing GPAI Compliance

Wireframe & Walkthrough

Atlas operationalizes an eight-stage AI compliance framework spanning the full model lifecycle, from pre-training through deployment authorization and continuous monitoring. Each staged activity shown below maps to specific obligations under the EU AI Act and adjacent regimes — that regulatory mapping sits behind each activity and is accessible through its detail view. What the dashboard presents is the second-line compliance officer’s vantage point: which activities are complete, who owns each one in the first line, and what remains for second-line review and sign-off before a frontier model can be responsibly deployed.

Conceptual design. Workspace, model, and personnel are illustrative.



Technical Documentation

- Training Data Inventory

External Engagement & Assessment

- AI Office Engagement Log
- AISI Engagement Log
- Code of Practice Commitments
- Transparency & Downstream Disclosures
- Conformity Assessment & Audit Bodies
- Voluntary Commitments

Baseline GPAI Documentation	Transparency & Disclosure	AI Literacy & Training	External Engagement	Adjacent Regulatory Regimes	Internal Governance
0/4	0/4	1/3	0/6	0/10	0/6
1ST LINE CONTROL			2ND LINE		
			OWNER	REVIEW	SIGN-OFF
01 · PRE-TRAINING			3 of 3 complete		
✓	Model Architecture Design	AW Architecture WG	✓	—	
✓	Training Planning	TE Training Eng	✓	—	
✓	Dataset Selection	DG Data Governance	✓	✓	
02 · TRAINING			2 of 2 complete		
✓	Early Capability Evaluation	EV Evals	✓	—	
✓	Safety Hypothesis	AL Alignment	✓	—	
03 · PRE-DEPLOYMENT EVALUATION			2 of 6 complete		
✓	Phase 2 Capability Evaluation	EV Evals	✓	—	
✓	Adversarial Testing	RT Red Team	✓	✓	
○	Red-Teaming Evaluation	RT Red Team	⌚	—	
○	Safeguard Testing	TS Trust & Safety	⌚	⌚	
○	Deployment Simulations	DP Deployment	⌚	—	
○	AISI Pre-Deployment Evaluation	EX External Engmt	⌚	⌚	

20

GOVERNANCE

Quarterly presentation · Altair-3

JUN 03

! Serious-incident report

REPORTING

AI Office filing — SI-2026-04 · Legal

12 days

JUN 04

AI Office: scheduled filing

FILING

Technical documentation submission

13 days

JUN 18

AISI Pre-Deployment Eval cycle

EVALUATION

Window opens — 14d evaluation

27 days

JUL 02

Serious Incident reporting window

REPORTING

Quarterly reporting close

41 days

JUL 15

Code of Practice quarterly report

REPORTING

Signatory reporting cycle Q2

54 days

AUG 01

Internal eval cycle milestone

EVALUATION

Phase 3 capability re-eval

71 days

RECENT ACTIVITY

●

Phase 2 Capability evals — evidence packet sealed

Evidence Log entry EV-2026-0411 sealed by L. Schneider. Chain-of-custody hash recorded.

EVIDENCE

EV-2026-0411

2 hours ago

04 · PRE-DEPLOYMENT COMPLIANCE ASSESSMENT0 of 2 complete

Cybersecurity Assessment

CO CISO Office

Systemic-Risk Analysis

CR CRO Office

05 · DEPLOYMENT AUTHORIZATION0 of 5 complete

Model Report

LA Legal & Audit

Risk Assessment

CR CRO Office

Mitigation Summary

MW Mitigation WG

Proposed Deployment Controls

DP Deployment

Residual Risk Recommendation

CR CRO Office

06 · BASELINE GPAI DOCUMENTATION0 of 4 complete

Technical Documentation (Annex XI)

LA Legal & Audit

Downstream Provider Information Package

LA Legal & Audit

Copyright Policy

LA Legal & Audit

Training Data Summary (Public)

DG Data Governance

07 · TRANSPARENCY & DISCLOSURE0 of 4 complete

AI Interaction Disclosure Mechanisms

PP Product Policy

AI-Generated Content Marking

PP Product Policy

Incident IR-2026-091 root-cause analysis

Hallucination cluster identified by Evals. Investigation underway — not flagged for escalation.

INCIDENTIR-2026-0913 hours ago

AI Office filing — drafting closes Friday

Technical documentation packet locked for review. Awaiting CRO sign-off.

AI OFFICEAI0-FIL-2026-Q2

4 hours ago

Red-Teaming Evaluation review awaiting sign

2nd-line review pending from CRO Office. Sign-off not required for this control.

REVIEWDRI: M. Thompson

5 hours ago

Mitigation MIT-118 marked effective

Effectiveness re-test passed (97.4%). Residual risk reduced from High to Moderate.

CONTROLMIT-118Yesterday

Vendor risk assessment completed — Tessera Labs

Third-party fine-tuning provider re-scored. No findings. DPA renewed through 2027.

VENDORVR-2026-094Yesterday

AI Office query on training data provenance

Follow-up question received re: synthetic data lineage. Response drafted, awaiting Legal review.

AI OFFICEAI0-Q-2026-022

2 days ago

AISI access agreement — addendum 3 executed

Updated access agreement. All data processed

☐	Deepfake Labeling Mechanisms	PP Product Policy		
☐	Automated Decision-Making Notices	LA Legal & Audit		
08 · AI LITERACY & TRAINING		1 of 3 complete		
☒	Staff AI Literacy Program	PO People Ops		
☐	Role-Specific AI Training Records	PO People Ops		
☐	Training Effectiveness Assessment	PO People Ops		
09 · EXTERNAL ENGAGEMENT		0 of 6 complete		
☐	AI Office Engagement Log	EX External Engmt		
☐	National Competent Authority Engagement	EX External Engmt		
☐	Code of Practice Adherence Statement	LA Legal & Audit		
☐	AISI / Safety Institute Engagement	EX External Engmt		
☐	Serious Incident Reporting Procedures	TS Trust & Safety		
☐	Conformity Assessment Body Coordination	EX External Engmt		
10 · ADJACENT REGULATORY REGIMES		0 of 10 complete		
☐	GDPR Legal Basis Assessment (training data)	LA Legal & Audit		
☐	Data Protection Impact Assessment	LA Legal & Audit		

EXECUTED

Updated scope covers Altair-3 Opus and Sonnet pre-deployment evaluations.

AISI AISI-AGR-3 3 days ago

☐	Cross-Border Transfer Mechanisms	LA Legal & Audit		
☐	TDM Opt-Out Compliance Mechanisms	DG Data Governance		
☐	Training Data Licensing Documentation	LA Legal & Audit		
☐	Export Control Classification	LA Legal & Audit		
☐	Model Weight Distribution Controls	CO CISO Office		
☐	Voluntary Commitment Adherence (Hiroshima)	CF Compliance		
☐	Industry Coalition Commitments (FMF, PAI)	EX External Engmt		
☐	Cross-Jurisdictional Compliance Assessment	LA Legal & Audit		
11 · INTERNAL GOVERNANCE 0 of 6 complete				
☐	Quality Management System Documentation	CF Compliance		
☐	Frontier Model Safety Committee Sign-Off	FS Frontier Safety		
☐	Board / Governance Body Notification	XO Exec Office		
☐	Executive Risk Acceptance	XO Exec Office		
☐	Legal Sign-Off on Compliance Posture	LA Legal & Audit		
☐	Compliance Officer Pre-Deployment Attestation	CF Compliance		